





6



System Design

There are two ways of constructing a software design: One way is to make it so simple that there are obviously no deficiencies, and the other way is to make it so complicated that there are no obvious deficiencies.
—C.A.R. Hoare






System design is the transformation of the analysis model into a system design model. During system design, developers define the design goals of the project and decompose the system into smaller subsystems that can be realized by individual teams. Developers also select strategies for building the system, such as the hardware/software platform on which the system will run, the persistent data management strategy, the global control flow, the access control policy, and the handling of boundary conditions. The result of system design is a model that includes a clear description of each of these strategies, a subsystem decomposition, and a UML deployment diagram representing the hardware/software mapping of the system.



System design is not algorithmic. Professionals and academics have, however, developed pattern solutions to common problems and defined notations for representing software architectures. In this chapter, we first present these building blocks and then discuss the design activities that have impact on these buildings blocks. In particular, system design includes:

- the definition of design goals
- the decomposition of the system into subsystems
- the selection of off-the-shelf and legacy components
- the mapping of subsystem to hardware
- the selection of a persistent data management infrastructure
- the selection of an access control policy
- the selection of a global control flow mechanism
- the handling of boundary conditions

We conclude this chapter by describing management issues that are related to system design.







6.1 Introduction: A floor plan example

System design, object design, and implementation constitute the construction of the system. During these three activities, developers fill the gap between the system specification, produced during requirements elicitation and analysis, and the system that is delivered to the users. System design is the first step in this process and focuses on decomposing the system into manageable parts. During requirements elicitation and analysis, we concentrated on the purpose and the functionality of the system. During system design, we focus on the processes, data structures, and software and hardware components necessary to implement it. The challenge of system design is that many conflicting criteria and constraints need to be met when decomposing the system.

Consider, for example, the task of designing a residential house. After agreeing with the client on the number of rooms and floors, the size of the living area, and the location of the house, the architect must design the floor plan, that is, where the walls, doors, and windows should be located. He must do so according to a number of functional requirements: The kitchen should be close to the dining room and the garage, the bathroom should be close to the bedrooms, and so on. The architect can also rely on a number of standards when establishing the dimensions of each room and the location of the door: Kitchen cabinets come in fixed increments and beds come in standard sizes. Note, however, that the architect does not need to know the exact contents of each room and the layout of the furniture; on the contrary, these decisions should be delayed and left to the client.




Figure 6-1 shows three successive revisions to a floor plan for a residential house. We set out to satisfy the following constraints:


- This house should have two bedrooms, a study, a kitchen, and a living room area.
- The overall distance the occupants walk every day should be minimized.

- The use of daylight should be maximized.


To satisfy the above constraints, we assume that most of the walking will be done between the entrance door and the kitchen, when groceries are unloaded from the car, and between the kitchen and the living/dining area, when dishes are carried before and after the meals. The next walking path to minimize is the path from the bedrooms to the bathrooms. We assume that the occupants of the house will spend most of their time in the dining area and in the master bedroom.

In the first version of our floor plan (at the top of Figure 6-1), we find that the dining room is too far from the kitchen. To address this problem, we exchange it with bedroom 2 (see gray arrows in Figure 6-1). This also has the advantage of moving the living room to the south wall of the house. In the second revision, we find that the kitchen and the stairs are too far from the entrance door. To address this problem, we move the entrance door to the north wall. This allows us to rearrange bedroom 2 and move the bathroom closer to both bedrooms. The living area increased, and we satisfied all constraints we originally set out.





Initial version



Second version

Third version









Figure 6-1 Example of floor plan design. Three successive versions show how we minimize walking distance and take advantage of sunlight.





At this point, we can position the doors and the windows of each room to meet localized requirements. Once this is done, we have completed the design of the floor, without detailed knowledge of the layout of each individual room. Plans for plumbing, electrical lines, and heating ducts can proceed.

The design of a floor plan in architecture is similar to system design in software engineering. The whole is divided into simpler components (i.e., rooms, subsystems) and interfaces (i.e., doors, services), while taking into account nonfunctional (i.e., living area, response time) and functional (i.e., number of bedrooms, use cases) requirements. System design impacts implementation activities (i.e., the kitchen layout, the coding complexity of individual subsystems) and results in costly rework if changed later (i.e., moving walls, changing subsystem interfaces). The design of individual components is delayed until later.

In Section 6.2, we provide you with a bird's eye view of system design and its relationship to analysis. In Section 6.3, we describe the concept of subsystems and subsystem decomposition. In Section 6.4, we describe system design activities and use an example to illustrate how these building blocks can be used together. In Section 6.5, we describe management issues related to system design.

6.2 An overview of system design

Analysis results in the requirements model described by the following products:


- a set of *nonfunctional requirements* and *constraints*, such as maximum response time, minimum throughput, reliability, operating system platform, and so on



- a *use case model*, describing the functionality of the system from the actors' point of view
- an *object model*, describing the entities manipulated by the system
- a *sequence diagram* for each use case, showing the sequence of interactions among objects participating in the use case

The analysis model describes the system completely from the actors' point of view and serves as the basis of communication between the client and the developers. The analysis model, however, does not contain information about the internal structure of the system, its hardware configuration, or, more generally, how the system should be realized. System design is the first step in this direction. System design results in the following products:

- list of *design goals*, describing the qualities of the system that developers should optimize
- *software architecture*, describing the subsystem decomposition in terms of subsystem responsibilities, dependencies among subsystems, subsystem mapping to hardware, and major policy decisions such as control flow, access control, and data storage


The design goals are derived from the nonfunctional requirements. Design goals guide the decisions to be made by the developers especially when trade-offs are needed. The subsystem decomposition constitutes the bulk of system design. Developers divide the system into manageable pieces to deal with complexity: Each subsystem is assigned to a team and realized independently. In order for this to be possible, though, developers need to address system-wide issues when decomposing the system. In particular, they need to address the following issues:





- *Hardware/software mapping:* What is the hardware configuration of the system? Which node is responsible for which functionality? How is communication between nodes realized? Which services are realized using existing software components? How are these components encapsulated? Addressing hardware/software mapping issues often leads to the definition of *additional subsystems* dedicated to moving data from one node to another, dealing with concurrency, and reliability issues. Off-the-shelf components enable developers to realize complex services more economically. User interface packages and database management systems are prime examples of off-the-shelf components. Components, however, should be encapsulated to minimize dependencies on a particular component: A competing vendor may come up with a better component.
- *Data management:* Which data need to be persistent? Where should persistent data be stored? How are they accessed? Persistent data represents a bottleneck in the system on many different fronts: Most functionality in system is concerned with creating or manipulating persistent data. For this reason, access to the data should be fast and reliable. If retrieving data is slow, the whole system will be slow. If data corruption is likely, complete system failure is likely. These issues need to be addressed consistently at the system level. Often, this leads to the selection of a database management system and of an *additional subsystem* dedicated to the management of persistent data.
- *Access control:* Who can access which data? Can access control change dynamically? How is access control specified and realized? Access control and security are system-wide issues. The access control must be consistent across the system; in other words, the policy used to specify who can and cannot access certain data should be the same *across all subsystems*.
- *Control flow:* How does the system sequence operations? Is the system event driven? Can it handle more than one user interaction at a time? The choice of control flow has an impact on the interfaces of subsystems. If an event-driven control flow is selected, subsystems will provide event handlers. If threads are selected, subsystems need to guarantee mutual exclusion in critical sections.
- *Boundary conditions:* How is the system initialized? How is it shut down? How are exceptional cases detected and handled? System initialization and shutdown often represent the larger part of the complexity of a system, especially in a distributed environment. Initialization, shutdown, and exception handling have an impact on the interface of *all subsystems*.

Figure 6-2 depicts the activities of system design. Each activity addresses one of the issues we described above. Addressing any one of these issues can lead to changes in the subsystem decomposition and to raising new issues. As you will see when we describe each of these activities, system design is a highly iterative activity, constantly leading to the identification of new subsystems, the modification of existing subsystems, and system-wide revisions that impact all subsystems. But first, let us describe the concept of subsystem in more detail.






Figure 6-2 The activities of system design (UML activity diagram).

6.3 System design concepts

In this section, we describe subsystem decompositions and their properties in more detail. First, we define the concept of **subsystem** and its relationship to classes (Section 6.3.1). Next, we look at the interface of subsystems (Section 6.3.2): subsystems provide **services** to other subsystems. A service is a set of related operations that share a common purpose. During system design, we define the subsystems in terms of the services they provide. Later, during object design, we define the subsystem interface in terms of the operations they provide. Next, we look at two properties of subsystems: **coupling** and **coherence** (Section 6.3.3). Coupling measures the dependencies between two subsystems, whereas coherence measures the dependencies among classes within a subsystem. Ideal subsystem decomposition should minimize coupling and maximize coherence. Next, we look at **layering** and **partitioning**, two techniques for relating subsystems to each other (Section 6.3.4). Layering allows a system to be organized as a hierarchy of subsystems, each providing higher level services to the subsystem above it by using lower level services from the subsystems below it. Partitioning organizes subsystems as peers



that mutually provide different services to each other. In Section 6.3.5, we describe a number of typical software architectures that are found in practice. Finally, in Section 6.3.6, we describe UML deployment diagrams, which we use to represent the mapping between software subsystems and hardware components.

6.3.1 Subsystems and classes

In Chapter 2, *Modeling with UML*, we introduced the distinction between application domain and solution domain. In order to reduce the complexity of the application domain, we identified smaller parts called classes and organized them into packages. Similarly, to reduce the complexity of the solution domain, we decompose a system into simpler parts, called subsystems, which are made of a number of solution domain classes. In the case of complex subsystems, we recursively apply this principle and decompose a subsystem into simpler subsystems (see Figure 6-3).





Figure 6-3 Subsystem decomposition (UML class diagram).

For example, the accident management system we previously described can be decomposed into a DispatcherInterfacesubsystem, implementing the user interface for the Dispatcher; a FieldOfficerInterfacesubsystem, implementing the user interface for the FieldOfficer; an IncidentManagementsubsystem, implementing the creation, modification, and storage of Incidents; and a Notificationsubsystem, implementing the communication between FieldOfficerterminals and Dispatcherstations. This subsystem decomposition is depicted in Figure 6-4 using UML packages.

Several programming languages (e.g., Java and Modula-2) provide constructs for modeling subsystems (packages in Java, modules in Modula-2). In other languages, such as C or C++, subsystems are not explicitly modeled, in which case, developers use conventions for grouping classes (e.g., a subsystem can be represented as a directory containing all the files implementing the subsystem). Whether or not subsystems are explicitly represented in the programming language, developers need to carefully document the subsystem decomposition as subsystems are usually realized by different teams.








Figure 6-4 Subsystem decomposition for an accident management system (UML class diagram, collapsed view). Subsystems are shown as UML packages. Dashed arrows indicate dependencies between subsystems.



6.3.2 Services and subsystem interfaces

A subsystem is characterized by the **services** it provides to other subsystems. A service is a set of related operations that share a common purpose. A subsystem providing a notification service, for example, defines operations to send notices, look up notification channels, and subscribe and unsubscribe to a channel.




The set of operations of a subsystem that are available to other subsystems form the **subsystem interface**. The subsystem interface, also referred to as the application programmer interface (API), includes the name of the operations, their parameters, their types, and their return values. System design focuses on defining the services provided by each subsystem, that is, enumerating the operations, their parameters, and their high-level behavior. Object design will focus on defining the subsystem interfaces, that is, the type of the parameters and the return value of each operation.


The definition of a subsystem in terms of the services it provides helps us focus on its interface as opposed to its implementation. A good subsystem interface should provide as little information about its implementation. This allows us to minimize the impact of change when we revise the implementation of a subsystem. More generally, we want to minimize the impact of change by minimizing the dependencies among subsystems.

6.3.3 Coupling and coherence

Coupling is the strength of dependencies between two subsystems. If two subsystems are loosely coupled, they are relatively independent, and thus, modifications to one of the subsystems will have little impact on the other. If two subsystems are strongly coupled, modifications to one

subsystem is likely to have impact on the other. A desirable property of a subsystem decomposition is that subsystems are as loosely coupled as possible. This minimizes the impact that errors or future changes have on the correct operation of the system.







Consider, for example, a compiler in which a parse tree is produced by the syntax analysis subsystem and handed over to the semantic analysis subsystem. Both subsystems access and modify the parse tree. An efficient way for sharing large amounts of data is to allow both subsystems to access the nodes of the tree via attributes. This introduces, however, a tight coupling: Both subsystems need to know the exact structure of the parse tree and its invariants. Figures 6-5 and 6-6 show the effect of changing the parse tree data structure for two cases: The left columns show the class interfaces when attributes are used for sharing data; the right columns show the class interfaces when operations are used. Because both the syntax analyzer and the semantic analyzer depend on these classes, both subsystems would need to be modified and retested in the case depicted by the left column. In general, sharing data via attributes increases coupling and should be avoided.

Coherence is the strength of dependencies within a subsystem. If a subsystem contains many objects that are related to each other and perform similar tasks, its coherence is high. If a subsystem contains a number of unrelated objects, its coherence is low. A desirable property of a subsystem decomposition is that it leads to subsystems with high coherence.

For example, consider a decision tracking system for recording design problems, discussions, alternative evaluations, decisions, and their implementation in terms of tasks (see Figure 6-7).



Binary tree representation




Sharing through attributes

```
class OpNode {  
  ArgNode left;  
  
  ArgNode right;  
  String name;  
}  
  
class ArgNode {  
  String name;  
}
```

Sharing through operations

```
class OpNode {  
  Enumeration  
  getArguments();  
  String getName();  
}  
  
class ArgNode {  
  String getName();  
}
```

Figure 6-5 Example of coupling reduction (UML object diagram and Java declarations). This figure shows a parse tree for the expression “a + b + c”. The left column shows the interface of the OpNode class with sharing through attributes. The right column shows the interface of OpNode with sharing through operations. Figure 6-6 shows the changes for each case when a linked list is selected instead.



Linked list representation

Sharing through attributes

```
class OpNode {  
    ArgNode first; ArgNode left; ArgNode right;  
    String name; }  
class ArgNode {  
    String name;  
    ArgNode next;  
}
```

Sharing through operations

```
class OpNode { Enumeration getArguments(); String getName();  
}  
  
class ArgNode { String getName(); }
```

Figure 6-6 Example of coupling reduction (UML object diagram and Java declarations). The figure shows the impact of changing the parse tree representation of Figure 6-5 to a linked list. In the left column, with sharing through attributes, four attributes need to change (changes indicated in *italics*). In the right column, with sharing through operations, the interface remains unchanged.

DecisionSubsystem






Figure 6-7 Decision tracking system (UML class diagram). The DecisionSubsystem has a low coherence: The classes Criterion, Alternative, and DesignProblem have no relationships with Subtask, ActionItem, and Task.





DesignProblem and Alternative represent the exploration of the design space: We formulate the system in terms of a number of DesignProblems and document each Alternative they explore. The Criterion class represents the qualities in which we are interested. Once we assessed the explored Alternatives against desirable Criteria, we take Decisions and implement them in terms of Tasks. Tasks are recursively decomposed into Subtasks small enough to be assigned to individual developers. We call atomic tasks ActionItems.

The decision tracking system is small enough that we could lump all these classes into one subsystem called DecisionSubsystem (see Figure 6-7). However, we observe that the class model can be partitioned into two subgraphs. One, called the RationaleSubsystem

, contains the classes DesignProblem, Alternative, Criterion, and Decision. The other, called the PlanningSubsystem, contains Task, Subtask, and ActionItem (see Figure 6-8). Both subsystems have a better coherence than the original DecisionSubsystem. Moreover, the resulting subsystems are smaller than the original subsystem: We reduced complexity. The coupling between the new subsystems is relatively low, as there is only one association between the two subsystems.









Figure 6-8 Alternative subsystem decomposition for the decision tracking system of Figure 6-7 (UML class diagram). The coherence of the RationaleSubsystem and the PlanningSubsystem is higher than the coherence of the original DecisionSubsystem. Note also that we also reduced the complexity by decomposing the system into smaller subsystems.







In general, there is a trade-off between coherence and coupling. We can always increase coherence by decomposing the system into smaller subsystems. However, this also increases coupling as the number of interfaces increases. A good heuristic is that developers can deal with 7 ± 2 concepts at any one level of abstraction. If there are more than nine subsystems at any given level of abstraction or if there is a subsystem providing more than nine services, you should consider a revision of the decomposition. By the same token, the number of layers should not be more than 7 ± 2 . In fact, many good systems design can be done with just three layers.

6.3.4 Layers and partitions

The goal of system design is to manage complexity by dividing the system into smaller, manageable pieces. This can be done by a divide-and-conquer approach, where we recursively divide parts until they are simple enough to be handled by one person or one team. Applying this approach systematically leads to a hierarchical decomposition in which each subsystem, or **layer**, provides higher level services, using services provided from lower level subsystems (see Figure 6-9). Each layer can only depend on lower level layers and has no knowledge of the layers above it. In a **closed architecture**, each layer can only depend on the layers immediately below it. In an **open architecture**, a layer can also access layers at deeper levels.





Layer 1 (Top) Layer 2 Layer 3 (Bottom)

Figure 6-9 Subsystem decomposition of a system into three layers (UML object diagram). A subset from a layered decomposition that includes at least one subsystem from each layer is called a vertical slice. For example, the subsystems A, B, and E constitute a vertical slice, whereas the subsystems D and G do not.

An example of a closed architecture is the Reference Model of Open Systems Interconnection (in short, the OSI model), which is composed of seven layers [Day & Zimmermann, 1983]. Each layer is responsible for performing a well-defined function. In addition, each layer provides its services by using services of the layer below.

The Physical layer represents the hardware interface to the network. It is responsible for transmitting bits over a communication channels. The Data Link layer is responsible for transmitting data frames without error using the services of the Physical layer. The Network layer is responsible for transmitting and routing packets within a network. The Transport layer is responsible for ensuring that the data are reliably transmitted from end to end. The Transport








Figure 6-10 An example of closed architecture: the OSI model (UML class diagram). The OSI model decomposes network services into seven layers, each responsible for a different level of abstraction.



layer is the interface Unix programmers see when transmitting information over TCP/IP sockets between two processes. The Sessionlayer is responsible for the initialization of a connection, including authentication. The Presentationlayer performs data transformation services, such as byte swapping or encryption. The Applicationlayer is the system you are designing (unless you are building an operating system or protocol stack). The application layer can also consist of layered subsystems.

Until recently, only the four bottom layers of the OSI model were well standardized. Unix and many desktop operating systems, for example, provide interfaces to TCP/IP that implemented the Transport, Network, and Datalinklayers. The application developer still needed to fill the gap between the Transportlayer and the Applicationlayer. With the growing number of distributed applications, this gap motivated the development of middleware such as CORBA [OMG, 1995] and Java RMI [RMI, 1998]. CORBA and Java RMI allow to **Format Connection Message Packet Frame Bit**








Figure 6-11 An example of closed architecture (UML class diagram). CORBA enables the access of objects implemented in different languages on different hosts. CORBA effectively implements the Presentation and Session layers of the OSI stack.

access remote objects transparently, by sending messages to them as messages are sent to local objects, effectively implementing the Presentation and Session layers (see Figure 6-11).

An example of an open architecture is the Motif user interface toolkit for X11 [Nye & O'Reilly, 1992]. The lowest layer, Xlib, provides basic drawing facilities and defines the concept of window. Xt provides tools for manipulating user interface objects, called widgets, using services from Xlib. Motif is a widget library that provides a wide range of facilities, from buttons to geometry management. Motif is built on top of Xt but also accesses Xlib directly. Finally, an application using Motif, such as a window manager, can access all three layers. Motif has no knowledge of the window manager and Xt has no knowledge of Motif or of the application. Many other user interface toolkits for X11 have open architectures. The openness of the architecture allows developers to bypass the higher level layers in case of performance bottleneck (Figure 6-12). **Object Socket Wire**










Figure 6-12 An example of open architecture: the OSF/Motif library (UML class diagram, packages collapsed). `Xlib` provides low-level drawing facilities. `Xt` provides basic user interface widget management. `Motif` provides a large number of sophisticated widgets. The `Application` can access each of these layers independently.




Closed layered architectures have desirable properties: They lead to low coupling between subsystems, and subsystems can be integrated and tested incrementally. Each level, however, introduces a speed and storage overhead that may make it difficult to meet nonfunctional requirements. Also, adding functionality to the system in later revisions may prove difficult, especially when the additions were not anticipated. In practice, a system is rarely decomposed into more than three to five layers.

Another approach to dealing with complexity is to **partition** the system into peer subsystems, each responsible for a different class of services. For example, an onboard system for a car could be decomposed into a travel service, giving real-time directions to the driver, an individual preferences service, remembering a driver's seat position and favorite radio station, and vehicle service, keeping track of the car's gas consumption, repairs, and scheduled maintenance. Each subsystem depends loosely on each other but could often operate in isolation.

In general, a subsystem decomposition is the result of both partitioning and layering. We first partition the system into top-level subsystems, which are responsible for specific functionality or which run on a specific hardware node. Each of the resulting subsystems are, if complexity justifies it, decomposed into lower and lower level layers until they are simple enough to be implemented by a single developer. Each subsystem adds a certain processing overhead because of its interface with other systems. Excessive partitioning or layering can lead to increased complexity.





6.3.5 Software architecture

As the complexity of systems increases, the specification of the system decomposition is critical. It is difficult to modify or correct a weak decomposition once development has started, as most subsystem interfaces have to change. In recognition of the importance of this problem, the concept of **software architecture** has emerged. A software architecture includes the system decomposition, the global control flow, error-handling policies and intersubsystem communication protocols [Shaw & Garlan, 1996].

In this section, we describe a few sample architectures that can be used for different systems. This is by no means a systematic or thorough exposition of the subject. Rather, we aim to provide a few representative examples and refer the reader to the literature for more details.

Repository architecture

In the repository architecture (see Figure 6-13), subsystems access and modify data from a single data structure called the central **repository**. Subsystems are relatively independent and interact only through the central data structure. Control flow can be dictated either by the central repository (e.g., triggers on the data invoke peripheral systems) or by the subsystems (e.g., independent flow of control and synchronization through locks in the repository).








Figure 6-13 Repository architecture (UML class diagram). Every subsystem depends only on a central data structure called the repository. The repository in turn, has no knowledge of the other subsystems.

The repository architecture is typical for database management systems, such as a payroll system or a bank system. The central location of the data makes it easier to deal with concurrency and integrity issues between subsystems. Modern compilers and software development environments also follow a repository architecture (see Figure 6-14). The different subsystems of a compiler access and update a central parse tree and a symbol table. Debuggers and syntax editors access the symbol table as well.

The repository subsystem can also be used for implementing the global control flow. In the compiler example of Figure 6-14, each individual tool (e.g., the compiler, the debugger, and the editor) is invoked by the user. The repository only ensures that concurrent accesses are serialized. Conversely, the repository can be used to invoke the subsystems based on the state of the central data structure. These systems are called blackboard systems. The HEARSAY II Speech understanding system [Erman et al., 1980], one of the first blackboard systems, selected tools to invoke based on the current state of the blackboard.








Figure 6-14 An instance of the repository architecture (UML Class diagram). A modern compiler incrementally generates a parse tree and a symbol table that can be later used by debuggers and syntax editors.

Repository architectures are well suited for applications with constantly changing complex data processing tasks. Once a central repository is well defined, we can easily add new services in the form of additional subsystems. The main disadvantage of repository systems is that the central repository can quickly become a bottleneck, both from a performance aspect and a modifiability aspect. The coupling between each subsystem and the repository is high, thus making it difficult to change the repository without having an impact on all subsystems.

Model/View/Controller

In the Model/View/Controller (MVC) architecture (see Figure 6-15), subsystems are classified into three different types: **model subsystems** are responsible for maintaining domain knowledge, **view subsystems** are responsible for displaying it to the user, and **controller subsystems** are responsible for managing the sequence of interactions with the user. The model subsystems are developed such that they do not depend on any view or controller subsystem. Changes in their state are propagated to the view subsystem via a subscribe/notify protocol. The MVC architecture is a special case of repository architecture where Model implements the central data structure and control objects dictate the control flow.






Figure 6-15 Model/View/Controller architecture (UML class diagram). The Controller gathers input from the user and sends messages to the Model. The Model maintains the central data structure. The View(s) display the Model and is notified (via a subscribe/notify protocol) whenever the Model is changed.

For example, Figures 6-17 and 6-16 illustrate the sequence of events that occur in a MVC architecture. Figure 6-17 displays two views of a file system. The bottom window lists the content of the Comp-Based Software Engineering folder, including the file 9DesignPatterns2.ppt

. The top window displays information about this file. The name of the file 9DesignPatterns2.ppt appears in three places: in both windows and in the title of the top window. Assume now that we change the name of the file to 9DesignPatterns.ppt

. Figure 6-16 shows the sequence of events:

1. The InfoView and the FolderView both subscribe for changes to the File models they display (when they are created).
2. The user types the new name of the file.
3. The Controller, the object responsible for interacting with the user during file name changes, sends a request to the Model.
4. The Model changes the file name and notifies all subscribers of the change.
5. Both InfoView and FolderView are updated, so the user sees a consistent change.




Figure 6-16 Sequence of events in the Model/View/Control architecture (UML collaboration diagram).






Figure 6-17 An example of MVC architecture. The “model” is the filename 9DesignPatterns2.ppt. One “view” is a window titled Comp-Based Software Engineering, which displays the contents of a folder containing the file 9DesignPatterns2.ppt. The other “view” is window called 9DesignPatterns2.ppt Info, which displays information related to the file. If the file name is changed, both views are updated by the “controller”.







Figure 6-17 An example of MVC architecture. The “model” is the filename 9DesignPatterns2.ppt. One “view” is a window titled Comp-Based Software Engineering, which displays the contents of a folder containing the file 9DesignPatterns2.ppt. The other “view” is window called 9DesignPatterns2.ppt Info, which displays information related to the file. If the file name is changed, both views are updated by the “controller”.



The subscription and notification functionality associated with this sequence of events is usually realized with an Observer pattern (see Section A.7). The **Observer pattern** allows the Model and the View objects to be further decoupled by removing direct dependencies from the Model to the View. For more details, the reader is referred to [Gamma et al., 1994] and to Section A.7 on page 504.

The rationale between the separation of Model, View, and Controller is that user interfaces, i.e., the View and the Controller, are much more often subject to change than is domain knowledge, i.e., the Model. Moreover, by removing any dependency from the Model on the View with the subscription/notification protocol, changes in the views (user interface) do not have any effect on the model subsystems. In the example of Figure 6-17, we could add a Unixstyle shell view of the file system without having to modify the file system. We described a similar decomposition in Chapter 5, *Analysis*, when we identified entity, boundary, and control objects. This decomposition is also motivated by the same considerations about change.





MVC architectures are well suited for interactive systems, especially when multiple views of the same model are needed. MVC can be used for maintaining consistency across distributed data; however, it introduces the same performance bottleneck as for other repository architectures.

Client/server architecture

In the client/server architecture (Figure 6-18), a subsystem, the **server**, provides services to instances of other subsystems called the **clients**, which are responsible for interacting with the user. The request for a service is usually done via a remote procedure call mechanism or a common object broker (e.g., CORBA or Java RMI). Control flow in the clients and the servers is independent except for synchronization to manage requests or to receive results.





Figure 6-18 Client/server architecture (UML class diagram). Clients request services from one or more servers. The server has no knowledge of the client. The client/server architecture is a generalization of the repository architecture.

An information system with a central database is an example of a client/server architecture. The clients are responsible for receiving inputs from the user, performing range checks, and initiating database transactions once all necessary data are collected. The server is then responsible for performing the transaction and guaranteeing the integrity of the data. In this case, a client/server architecture is a special case of the repository architecture, where the central data structure is managed by a process. Client/server systems, however, are not restricted to a single server. On the World Wide Web, a single client can easily access data from thousands of different servers (Figure 6-19).








Figure 6-19 The World Wide Web as an instance of the client/server architecture (UML object diagram).





Client/server architectures are well suited for distributed systems that manage large amounts of data.

Peer-to-peer architecture

A peer-to-peer architecture (see Figure 6-20) is a generalization of the client/server architecture in which subsystems can act both as client or as servers, in the sense that each subsystem can request and provide services. The control flow within each subsystem is independent from the others except for synchronizations on requests.




Figure 6-20 Peer-to-peer architecture (UML class diagram). Peers can request services from and provide services to other peers.






Figure 6-21 An example of peer-to-peer architecture (UML collaboration diagram). The database server can both process requests from and send notifications to applications.

An example of a peer-to-peer architecture is a database which, on the one hand, accepts requests from the application and, on the other hand, sends notifications to the application whenever certain data are changed. Peer-to-peer systems are more difficult to design than client/ server systems are. They introduce the possibility of deadlocks and complicate the control flow.

Pipe and filter architecture

In the pipe and filter architecture (see Figure 6-22), subsystems process data received from a set of inputs and send results to other subsystems via a set of outputs. The subsystems are called **filters**, and the associations between the subsystems are called **pipes**. Each filter only knows the content and the format of the data received on the input pipes, not the filters that produced them. Each filter is executed concurrently and synchronization is done via the pipes. The pipe and filter architecture is modifiable: Filters can be substituted for others or reconfigured to achieve a different purpose.






Figure 6-22 Pipe and filter architecture (UML class diagram). A `Filter` can have many inputs and outputs. A `Pipe` connects one of the outputs of a `Filter` to one of the inputs of another `Filter`.

The best known example of pipe and filter architecture is the Unix shell. Most filters are written such that they read their input and write their results on standard pipes. This enables a Unix user to combine them in many different ways. Figure 6-23 shows an example made out of four filters. The output of `ps`(process status) is fed into `grep`(search for a pattern) to get rid of all the processes that are not owned by a specific user. The output of `grep`(i.e., the processes owned by the user) is then sorted by `sort` and then sent to `more`, which is a filter that displays its input to a terminal, one screen at a time.




Figure 6-23 An instance of the pipe and filter architecture (Unix command and UML activity diagram).


Pipe and filter architectures are suited for systems that apply transformations to streams of data without intervention by users. They are not suited for systems that require more-complex interactions between components, such as an information management system or an interactive system.

6.3.6 UML deployment diagrams

UML deployment diagrams are used to depict the relationship among run-time components and hardware nodes. Components are self-contained entities that provide services to other components or actors. A Web server, for example, is a component that provides services to Web browsers. A Web browser such as Netscape is a component that provides services to a user. A distributed system can be composed of many interacting run-time components.

In UML deployment diagrams, nodes are represented by boxes containing component icons. Dependencies between components are represented by dashed arrows. Figure 6-24 depicts an example of deployment diagram illustrating two Web browsers accessing a Web server. The

Web server in turns accesses data from a database server. We can see from the dependency graph that the Web browsers do not access directly the database at any time.







Figure 6-24 A UML deployment diagram representing the allocation of components to different nodes and the dependencies among components. Web browsers on PCs and Macs can access a WebServer that provides information from a Database.



The deployment diagram in Figure 6-24 focuses on the allocation of components to nodes and provides a high-level view of each component. Components can be refined to include information about the interfaces they provide and the classes they contain. Figure 6-25 illustrates the GET and POST interfaces of the WebServer component and its containing classes.








Figure 6-25 Refined view of the WebServer component (UML deployment diagram). The WebServer component provides two interfaces to browsers: A browser can either request the content of a file referred by a URL (GET) or post the content of a form (POST). The WebServer component contains five classes: URL, HttpRequest, DBQuery, File, and DBResult.





6.4 System design activities: From objects to subsystems

System design consists of transforming the analysis model into the design model that takes into account the nonfunctional requirements and constraints described in the problem statement and the requirements analysis document. In Section 6.3, we focused on subsystem decompositions and their properties. In this section, we describe the activities that are needed to ensure that a subsystem decomposition addresses all the nonfunctional requirements and prepares for taking into account any constraints during the implementation phase. We illustrate these activities with an example, MyTrip, a route planning system for car drivers, throughout this section. This will provide you with more concrete knowledge of system design concepts.

We start with the analysis model from MyTrip. We then:


- identify design goals from the nonfunctional requirements (Section 6.4.2)
- design an initial subsystem decomposition (Section 6.4.3)
- map subsystems to processors and components (Section 6.4.4)
- decide storage (Section 6.4.5)
- define access control policies (Section 6.4.6)
- select a control flow mechanism (Section 6.4.7)
- identify boundary conditions (Section 6.4.8)


In Section 6.4.9, we examine issues related to stabilizing the system design while anticipating change. Finally, in Section 6.4.10, we describe

how the system design model is reviewed. But first, we describe the analysis model we use as a starting point for the system design of MyTrip.

6.4.1 Starting point: Analysis model for a route planning system


Using MyTrip, a driver can plan a trip from a home computer by contacting a trip planning service on the Web (PlanTripin Figure 6-26). The trip is saved for later retrieval on the server. The trip planning service must support more than one driver.







<i>Use case name</i>	PlanTrip
<i>Entry condition</i>	1. The Driver activates her home computer and logs into the trip planning Web service.
<i>Flow of events</i>	2. Upon successful login, the Driver enters constraints for a trip as a sequence of destinations. 3. Based on a database of maps, the planning service computes the shortest way visiting the destinations in the specified order. The result is a sequence of segments binding a series of crossings and a list of directions. 4. The Driver can revise the trip by adding or removing destinations.
<i>Exit condition</i>	5. The Driver saves the planned trip by name in the planning service database for later retrieval.







The driver then goes to the car and starts the trip. While the onboard computer gives directions based on trip information from the planning service and her current position indicated by an onboard GPS system. (ExecuteTrip in Figure 6-27).



<i>Use case name</i>	ExecuteTrip
<i>Entry condition</i>	1. The Driver starts her car and logs into the onboard route assistant.
<i>Flow of events</i>	2. Upon successful login, the Driver specifies the planning service and the name of the trip to be executed. 3. The onboard route assistant obtains the list of destinations, directions, segments, and crossings from the planning service. 4. Given the current position, the onboard route assistant provides the next set of directions.
<i>Exit condition</i>	5. The Driver arrives to destination

We perform the analysis for the MyTrip system following the techniques outlined in Chapter 5, *Analysis*, and obtain the model in Figure 6-28.






Crossing	A Crossing is a geographical point where a driver can choose between several Segments.
Destination	A Destination represents a location where the driver wishes to go.
Direction	Given a Crossing and an adjacent Segment, a Direction describes in natural language terms how to steer the car onto the given Segment.
Location	A Location is the position of the car as known by the onboard GPS system or the number of turns of the wheels.
PlanningService	A PlanningService is a Web server that can supply a trip, linking a number of destinations in the form of a sequence of crossings and segments.
RouteAssistant	A RouteAssistant gives Directions to the driver, given the current Location and upcoming Crossing.
Segment	A Segment represents the road between two Crossings.

In addition, during requirements elicitation, our client specified the following nonfunctional requirements for MyTrip.

Nonfunctional requirements for MyTrip

1. MyTrip is in contact with the PlanningService via a wireless modem. It can be assumed that the wireless modem functions properly at the initial destination.
2. Once the trip has been started, MyTrip should give correct directions even if modem fails to maintain a connection with the PlanningService.
3. MyTrip should minimize connection time to reduce operation costs.
4. Replanning is possible only if the connection to the PlanningService is possible.
5. The PlanningService can support at least 50 different drivers and 1000 trips.



6.4.2 Identifying design goals


The definition of design goals is the first step of system design. It identifies the qualities that our system should focus on. Many design goals can be inferred from the nonfunctional requirements or from the application domain. Others will have to be elicited from the client. It is, however, necessary to state them explicitly such that every important design decision can be made consistently following the same set of criteria.

For example, in the light of the nonfunctional requirements for MyTrip described in Section 6.4.1, we identify *reliability* and *fault tolerance to connectivity loss* as design goals. We then identify *security* as a design goal, as numerous drivers will have access to the same trip planning server. We add *modifiability* as a design goal, as we want to provide the ability for drivers to select a trip planning service of their choice. The following box summarizes the design goals we identified.







Design goals for MyTrip

- **Reliability:** MyTrip should be reliable [generalization of nonfunctional requirement 2].
- **Fault Tolerance:** MyTrip should be fault tolerant to loss of connectivity with the routing service [rephrased nonfunctional requirement 2].
- **Security:** MyTrip should be secure, i.e., not allow other drivers or nonauthorized users to access another driver's trips [deduced from application domain].
- **Modifiability:** MyTrip should be modifiable to use different routing services [anticipation of change by developers].



In general, we can select design goals from a long list of highly desirable qualities. Tables 6-1–6-5 list a number of possible design criteria. These criteria are organized into five groups: *performance*, *dependability*, *cost*, *maintenance*, and *end user criteria*. Performance, dependability, and end user criteria are usually specified in the requirements or inferred from the application domain. Cost and maintenance criteria are dictated by the customer and the supplier.

Design criterion	Definition
Performance criteria (Table 6-1) include the speed and space requirements imposed on the system. Should the system be responsive or should it accomplish a maximum number of tasks? Is memory space available for speed optimizations or should memory be used sparingly?	
Response time	How soon is a user request acknowledged after the request has been issued?
Throughput	How many tasks can the system accomplish in a fixed period of time?
Memory	How much space is required for the






Table 6-2 Dependability criteria

Design criterion	Definition
Robustness	Ability to survive invalid user input
Reliability	Difference between specified and observed behavior.
Availability	Percentage of time system can be used to accomplish normal tasks.
Fault tolerance	Ability to operate under erroneous conditions.
Security	Ability to withstand malicious attacks
Safety	Ability to not endanger human lives, even in the presence of errors and



Dependability criteria (Table 6-2) determine how much effort should be expended in minimizing system crashes and their consequences. How often can the system crash? How available to the user should the system be? Are there safety issues associated with system crashes? Are there security risks associated with the system environment?

Cost criteria (Table 6-3) include the cost to develop the system, to deploy it, and to administer it. Note that cost criteria not only include design considerations but managerial ones as well. When the system is replacing an older one, the cost of ensuring backward compatibility or transitioning to the new system has to be taken into account. There are also trade-offs between different types of costs such as development cost, end user training cost, transition costs and maintenance costs. Maintaining backward compatibility with a previous system can add to the development cost while reducing the transition cost.



Table 6-3 Cost criteria

Design criterion	Definition
Development cost	Cost of developing the initial system
Deployment cost	Cost of installing install the system and training the users.
Upgrade cost	Cost of translating data from the previous system. This criteria results in backward compatibility requirements.
Maintenance	Cost required for bug fixes and





Maintenance criteria (Table 6-4) determine how difficult it is to change the system after deployment. How easily can new functionality be added? How easily can existing functions be revised? Can the system be adapted to a different application domain? How much effort will be required to port the system to a different platform? These criteria are harder to optimize and plan for, as it is seldom clear how long the system will be operational and how successful the project will be.



Table 6-4 Maintenance criteria

Design criterion	Definition
Extensibility	How easy is it to add the functionality or new classes of the system?
Modifiability	How easy is it to change the functionality of the system?
Adaptability	How easy is it to port the system to different application domains?
Portability	How easy is it to port the system to different platforms?
Readability	How easy is it to understand the system from reading the code?
Traceability of	How easy is it to map the code to

Table 6-5 End user criteria

Design criterion	Definition
Utility	How well does the system support the work of the user?
Usability	How easy is it for the user to use the system?

When defining design goals, only a small subset of these criteria can be simultaneously taken into account. It is, for example, unrealistic to develop software that is safe, secure, and cheap. Typically, developers need to prioritize design goals and trade them off against each other as well as against managerial goals as the project runs behind schedule or over budget. Table 6-6 lists several possible trade-offs.






Table 6-6 Examples of design goal trade-offs



Trade-off	Rationale
Space vs. speed	If the software does not meet response time or throughput requirements, more memory can be expended to speed up the software (e.g., caching, more redundancy, etc.). If the software does not meet memory space constraints, data can be compressed at the cost of speed.
Delivery time vs. functionality	If the development runs behind schedule, a project manager can deliver less functionality than specified and deliver on time, or deliver the full functionality at a later time. Contract software usually puts more emphasis on functionality, whereas off-the-shelf software projects puts more emphasis on delivery date.
Delivery time vs. quality	If the testing runs behind schedule, a project manager can deliver the software on time with known bugs (and possibly provide a later patch to fix any serious bugs) or to deliver the software later with more bugs fixed.
Delivery time vs. staffing	If development runs behind schedule, a project manager can add resources to the project in order to increase productivity. In most cases, this option is only available early in the project: Adding resources usually decreases productivity while new personnel is being trained or brought up to date. Note that adding resources will also raise the cost of development.

Managerial goals can be traded off against technical goals (e.g., delivery time vs. functionality). Once we have a clear idea of the design goals, we can proceed to design an initial subsystem decomposition.

6.4.3 Identifying subsystems

Finding subsystems during system design has many similarities to finding objects during analysis: It is a volatile activity driven by heuristics. As a result, the object identification techniques we described in Chapter 5, *Analysis*, such as Abbotts’s lexical rules, are applicable to subsystem identification. Moreover, subsystem decomposition is constantly revised whenever new issues are addressed: Subsystems are merged into one subsystem, a complex subsystem is split into parts, and some subsystems are added to take care of new functionality. The first iterations over the subsystem decomposition can introduce drastic changes in the system design model. These are often best handled through brainstorming.

The initial subsystem decomposition should be derived from the functional requirements. For example, in the MyTrip system, we identify two major groups of objects: those that are involved during the PlanTripuse cases and those that are involved during the ExecuteTripuse case. The Trip, Direction, Crossing, Segment, and Destinationclasses are shared between both use cases. This set of classes is tightly coupled as they are used as a whole to represent a Trip. We decide to assign them with PlanningServiceto the PlanningSubsystem, and the remainder of the classes are assigned to the RoutingSubsystem. This leads to only one association crossing subsystem boundaries. Note that this subsystem decomposition follows a repository architecture in which the PlanningSubsystemis responsible for the central data structure.







Figure 6-29 Initial subsystem decomposition for MyTrip (UML class diagram).

Another heuristic for subsystem identification is to keep functionally related objects together. A starting point is to take the use cases and assign the participating objects that have been identified in each of them to the subsystems. Some group of objects, as the Tripgroup in MyTrip, are shared and used for communicating information from one subsystem to another. We can either create a new subsystem to accommodate them or assign them to the subsystem that creates these objects.

Heuristics for grouping objects into subsystems

- Assign objects identified in one use case into the same subsystem
- Create a dedicated subsystem for objects used for moving data among subsystems
- Minimize the number of associations crossing subsystem boundaries
- All objects in the same subsystem should be functionally related



Encapsulating subsystems

Subsystem decomposition reduces the complexity of the solution domain by minimizing dependencies among classes. The **Facade pattern** [Gamma et al., 1994] allows us to further reduce dependencies between classes by encapsulating a subsystem with a simple, unified interface. For example, in Figure 6-30, the Compiler class is a Facade hiding the classes CodeGenerator, Optimizer, ParseNode, Parser, and Lexer. The Facade provides access only to the public services offered by the subsystem and hides all other details, effectively reducing coupling between subsystems.






Figure 6-30 An example of Facade pattern (UML class diagram).

Subsystems identified during the initial subsystem decomposition often result from grouping several functionally related classes. These subsystems are good candidates for the Facade patterns and should be encapsulated under one class.

6.4.4 Mapping subsystems to processors and components

Selecting a hardware configuration and a platform

Many systems run on more than one computer and depend on access to an intranet or to the Internet. The use of multiple computers can address high-performance needs or to interconnect multiple distributed users. Consequently, we need to carefully examine the allocation of subsystems to computers and the design of the infrastructure for supporting communication between subsystems. These computers are modeled as nodes in UML deployment diagrams. Nodes can either represent a specific instances (e.g., myMac) or a class of computers (e.g., WebServer).¹ Because the hardware mapping activity has significant impact on the performance and complexity of the system, we perform it early in system design.





Select a hardware configuration also includes selecting a virtual machine onto which the system should be built. The virtual machine includes the operating system and any software components that are needed, such as a database management system or a communication package. The selection of a virtual machine reduces the distance between the system and the hardware platform on which it will run. The more functionality the components provide, the less work involved. The selection of the virtual machine, however, may be constrained by the client who acquires hardware before the start of the project. The selection of a virtual machine may also be constrained by cost considerations: In some cases, it is difficult to estimate whether building a component costs more than buying it.

In MyTrip, we deduce from the requirements that PlanningSubsystem and RoutingSubsystem run on two different nodes: The former is a Web-based service on an Internet host while the second runs on the onboard computer. Figure 6-31 illustrates the hardware allocation for MyTrip with two nodes called :OnBoardComputer and :WebServer.




Figure 6-31 Allocation of MyTrip subsystems to hardware (UML deployment diagram). RoutingSubsystem runs on the OnBoardComputer while PlanningSubsystem runs on a WebServer.


We select a Unix machine as the virtual machine for the :WebServer and the Web browsers Netscape and Internet Explorer as the virtual machines for the :OnBoardComputer.

Allocating objects and subsystems to nodes

Once the hardware configuration has been defined and the virtual machines selected, objects and subsystems are assigned to nodes. This often triggers the identification of new objects and subsystems for transporting data among the nodes.

In the MyTrip system, both RoutingSubsystem and PlanningSubsystem share the objects Trip, Destination, Crossing, Segment, and Direction. Instances of these classes need to communicate via a wireless modem using some communication protocol. We create a new

1. These two cases are distinguished with the usual UML naming convention: underlined names for instances and nonunderlined names for classes.



Syst


subsystem to support this communication: CommunicationSubsystem

, a subsystem located on both nodes for managing the communication

between the two.

We also notice that only segments constituting the planned trip are stored in RoutingSubsystem. Adjacent segments not part of the trip are stored only in the PlanningSubsystem. To take this into account, we need objects in the RoutingSubsystem that can act as a surrogates to Segments and Trips in the PlanningSubsystem. An object that acts on the behalf of another one is called a proxy. We therefore create two new classes, called SegmentProxy and TripProxy, and make them part of the RoutingSubsystem. These proxies are examples of the **Proxy** design pattern [Gamma et al., 1994].

In case of replanning by the driver, this class will transparently request the CommunicationSubsystem to retrieve the information associated with its corresponding Segments on the PlanningSubsystem. Finally, the CommunicationSubsystem is used for transferring a complete trip from PlanningSubsystem to RouteAssistant. The revised design model and the additional class descriptions are depicted in Figure 6-32.



CommunicationSubsystem The CommunicationSubsystem is responsible for transporting objects from the PlanningSubsystem to the RoutingSubsystem.

Connection A Connection represents an active link between the PlanningSubsystem and the RoutingSubsystem. A Connection object handles exceptional cases associated with loss of network services.

Message A Message represents a Trip and its related Destinations, Segments, Crossings, and Directions, encoded for transport.

Figure 6-32 Revised design model for MyTrip (UML Class diagram, associations omitted for clarity).



In general, allocating subsystems to hardware nodes enables us to distribute functionality and processing power where it is most needed. Unfortunately, it also introduces issues related to storing, transferring, replicating, and synchronizing data among subsystems. For this reason, developers also select the components they will use for developing the system.

Encapsulating components

As the complexity of systems increases and the time to market shortens, developers have strong incentives to reuse code and to rely on vendor-supplied components. Interactive systems, for example, are now rarely built from scratch but rather are developed with user interface toolkits that provide a wide range of dialogs, windows, buttons, or other standard interface objects. Other projects focus on redoing only part of an existing system. For example, corporate information systems, costly to design and build, need to be updated to new client hardware. Often, only the client side of the system is upgraded to new technology and the backend of the system left untouched.² Whether dealing with off-the-shelf component or legacy code, developers have to deal with existing code, which they cannot modify and which has not been designed to be integrated into their system.

We can deal with existing components such as code by encapsulating them. This approach has the advantage of decoupling the system from the encapsulated code, thus minimizing the impact of existing software on the design. When the encapsulated code is written in the same language as the new system, this can be done using an Adapter pattern.

The **Adapter pattern** (Figure 6-33) is used to convert the interface of an existing piece of code into an interface, called the *NewInterface*, that a calling subsystem expects. An *Adaptor* class, also called a *wrapper*, is introduced to provide the glue between *NewInterface* and *LegacySystem*. For example, assume the client is the static `sort()` method of the Java *Array* class (Figure 6-34). This method expects two arguments *a*, an *Array* of objects, and *c*, a






Figure 6-33 Adapter pattern (UML class diagram). The Adapter pattern is used to provide a different interface (*New Interface*) to an existing component (*Legacy System*).

2. Such projects are called interface engineering projects (see Chapter 4, *Requirements Elicitation*).





```
/* Existing target interface */
interface Comparator { int compare(Object o1, Object o2); /* ... */
}

// Existing client
class Array { static void sort(Object [] a, Comparator c); /* ... */
}

/* Existing adaptee class */
class MyString extends String { boolean equals(Object o); boolean greaterThan(MyString s); /* ... */
}

/* New adaptor class */
class MyStringComparator implements Comparator {
    /* ... */
    int compare(Object o1, Object o2) { int result; if (o1.greaterThan(o2)) {
        result = 1
    } else if (o1.equals(o2)) { result = 0;
    } else {
        result = -1; } return result; }
}
```








Figure 6-34 Adapterpattern example (Java). The static sort()method on Arrays takes two arguments: an arrays of Objects to be sorted and a Comparatordefining the relative order of the elements. To sort an array of MyStrings, we need to define a comparator called MyStringComparator with the proper interface. MyStringComparator is an Adaptor.

Comparatorobject, which provides a compare()method, defining the relative order between elements. Assume we are interested in sorting strings of the class MyString, which defines the greaterThan()and an equals()methods. To sort an Arrayof MyStrings, we need to define a new comparator, MyStringComparatorcompare() method using

, which provides a greaterThan()and equals(). MyStringComparatoris an Adaptorclass.³

3. When designing a new system, adapters are seldom necessary, as the interface of new classes can be defined such that they comply with the necessary interfaces.





When encapsulating legacy code that is written in a different language than the system under development, we need to deal with language differences. Although integrating code from two different compiled languages can be done, it can present major problems, especially when one or both of the languages are object-oriented and implement different message dispatching semantics. This motivated standards such as CORBA, which defines protocols for allowing the interoperability of distributed objects written in different languages. In the case of client/server architectures, other solutions include developing wrappers around communication protocols between processes.

Protocols for interprocess communication (e.g., pipes and sockets) are usually provided by the operating system and, thus, are language independent. In the cases of CORBA and interprocess communication, the cost of invoking a service becomes much higher than the cost of sending messages among objects in the same process. You need to carefully evaluate the impact on performance by wrappers around legacy code when response time and other performance design goals have been selected.




Technology decisions become obsolete quickly. The system you are building is likely to survive many platforms and will be ported and upgrade several times during maintenance. These tasks, when performed during maintenance, are usually costly, because a large amount of the design information is lost. Why was the original hardware configuration used? Which features of the database management system does this system rely on? Developers can preserve this information by documenting the design rationale of their system, including hardware and component decisions. We describe techniques for doing this in Chapter 8, *Rationale Management*.

6.4.5 Defining persistent data stores

Persistent data outlive a single execution of the system. For example, an author may save his work into a file when using a word processor. The file can then be reopened several days or weeks later. The word processor need not run for the file to exist. Similarly, information related to employees, their employment status, and their paychecks live in a database management system. This allows all the programs that operate on employee data to do so consistently. Moreover, storing data in a database enables the system to perform complex queries on a large data set (e.g., the records of several thousands of employees).

Where and how data is stored in the system impacts the system decomposition. In some cases, for example, in a repository architecture (see Section 6.3.5), a subsystem can be completely dedicated to the storage of data. The selection of a specific database management system can also have implications on the overall control strategy and concurrency management.

For example, in MyTrip, we decide to store the current Tripin a file on a small removable disk in order to allow the recovery of the Tripin case the driver shuts off the car before reaching the final Destination. Using a file is the simplest and most efficient solution in this case, given that the RoutingSubsystemwill only store complete trips to the file before shutdown and load the file at start-up. In the PlanningSubsystem, however, the trips will be stored in a database.









Figure 6-35 Subsystem decomposition of MyTrip after deciding on the issue of data stores (UML class diagram, packages collapsed for clarity).

In general, we first need to identify which objects need to be persistent. The persistency of objects is directly inferred from the application domain. In MyTrip, only Trips and their related classes need to be stored. The location of the car, for example, need not be persistent because it needs to be recalculated constantly. Then, we need to decide how these objects should be stored (e.g., file, relational database, or object database). The decision for the storage management is more complex and is usually dictated by nonfunctional requirements: Should the objects be retrieved quickly? Is there a need for complex queries? Do objects take a lot of space (e.g., are there images to store)? Database management systems provide mechanisms for concurrency control and efficient queries over large datasets.





There are currently three realistic options for storage management:

- **Flat files.** Files are the storage abstractions provided by operating systems. The application stores its data as a sequence of bytes and defines how and when data should be retrieved. The file abstraction is relatively low level and enables the application to perform a variety of size and speed optimizations. Files, however, require the application to take care of many issues, such as concurrent accesses, and loss of data in case of system crash.
- **Relational database.** A relational database provides an abstraction of data that is higher than flat files. Data are stored in tables that comply with a predefined type called *schema*. Each column in the table represents an attribute. Each row represents a data item as a tuple of attribute values. Several tuples in different tables are used to represent the attributes of an individual object. Relational databases have been used for awhile and are a mature technology. The use of a relational database introduces a high cost and, often, a performance bottleneck.
- **Object-oriented database.** An object-oriented database provides services similar to a relational database. Unlike a relational database, it stores data as objects and associations. In addition to providing a higher level of abstraction (and thus reducing the need to translate between objects and storage entities), object-oriented databases provide developers with inheritance and abstract datatypes. Object-oriented databases are usually slower than relational databases for typical queries.



The following box summarizes the trade-offs when selecting storage management system.




Trade-off between files and databases

When should you choose a file?

- Voluminous data (e.g., images)
- Temporary data (e.g., core file)
- Low information density (e.g., archival files, history logs)
- Concurrent accesses
- Access at fine levels of details
- Multiple platforms
- Multiple applications over the same data
- Complex queries over attributes
- Large dataset
- Extensive use of associations to retrieve data
- Medium-sized data set
- Irregular associations among objects





Encapsulating data stores

Once we select a storage mechanism (say, a relational database), we can encapsulate it into a subsystem and define a high-level interface that is vendor independent. For example, the **Bridge pattern** (see Figure 6-36 and [Gamma et al., 1994]) allows the interface and the implementation of a class to be decoupled. This allows the substitution of different implementations of a given class, sometimes even at run-time. The *Abstraction* class defines the interface visible to the client. The *Implementor* is an abstract class that defines the lower level methods available to *Abstraction*. An *Abstraction* instance maintains a reference to its corresponding *Implementor* instance. *Abstraction* and *Implementor* can be refined independently.




Figure 6-36 Bridge pattern (UML class diagram).



Database connectivity standards such as ODBC [Microsoft, 1995] and JDBC [JDBC, 1998] provide such abstractions for relational databases (see ODBC Bridge pattern in Figure 6-37). Note, however, that even if most relational databases provide similar services, providing such an abstraction reduces performance. The design goals we defined at the beginning of the system design phase help us trade off performance and modifiability.







Figure 6-37 Bridge pattern for abstracting database vendors (UML class diagram). Removing the dependency from database vendors from the systems provides more flexibility.





6.4.6 Defining access control

In multiuser systems, different actors have access to different functionality and data. For example, an everyday actor may only access the data it creates, whereas a system administrator actor may have unlimited access to system data and to other users' data. During analysis, we modeled these distinctions by associating different use cases to different actors. During system design, we model access by examining the object model, by determining which objects are shared among actors, and by defining how actors can control access. Depending on the security requirements on the system, we also define how actors are authenticated to the system (i.e., how actors prove to the system who they are) and how selected data in the system should be encrypted.



Table 6-7 Revisions to the design model stemming from the decision to authenticate Drivers and encrypt communication traffic. The text added to the model is in *italics*.


CommunicationSubsystem	The CommunicationSubsystem is responsible for transporting Trips from the PlanningSubsystem to the RoutingSubsystem. <i>The CommunicationSubsystem uses the Driver associated with the Trip being transported for selecting a key and encrypting the communication traffic.</i>
PlanningSubsystem	The PlanningSubsystem is responsible for constructing a Trip connecting a sequence of Destinations. The PlanningSubsystem is also responsible for responding to replan requests from RoutingSubsystem. <i>Prior to processing any requests, the PlanningSubsystem authenticates the Driver from the RoutingSubsystem. The authenticated Driver is used to determine which Trips can be sent to the corresponding RoutingSubsystem.</i>
Driver	<i>A Driver represents an authenticated user. It is used by the CommunicationSubsystem to remember keys associated with a user and by the PlanningSubsystem to associate Trips with users.</i>

For example, in MyTrip, storing maps and Trips for many drivers in the same database introduces security issues. We must ensure that Trips are sent only to the driver that created them. This is also consistent with the security design goal we defined in Section 6.4.2 for MyTrip. Consequently, we model a driver with the Driverclass and associate it with the Tripclass. The PlanningSubsystemalso becomes responsible for authenticating Drivers before sending Trips. Finally, we decide to encrypt the communication traffic between the RoutingSubsystem and the PlanningSubsystem. This will be done by the CommunicationSubsystem

. The descriptions for the Driverclass and the revised descriptions for the PlanningSubsystemand the CommunicationSubsystemare displayed in Table 6-7. The revisions to the design model are indicated in italics.

Defining access control for a multiuser system is usually more complex than in MyTrip. In general, we need to define for each actor which operations they can access on each shared object. For example, in a bank information system, a teller may credit or debit money from local accounts up to a predefined amount. If the transaction exceeds the predefined amount, a manager needs to approve the transaction. Moreover, managers and tellers can only access accounts in their own branch; that is, they cannot access accounts in other branches. Analysts, on the other hand, can access information across all branches of the corporation, but cannot post transactions on individual accounts.







We model access on classes with an access matrix. The rows of the matrix represents the actors of the system. The columns represent classes whose access we control. An entry (class, actor) in the access matrix is called an **access right** and lists the operations (e.g., postSmallDebit(), postLargeDebit(), getCustomerAddress()



, examineBalance()) that can be executed on instances of the classby the actor. Table 6-8 is an example of access matrix for our bank information system.

Table 6-8 Access matrix for a banking system. Tellers can only lookup local accounts, perform small transactions on accounts, and inquire balances. Managers can perform larger transactions and access account history in addition to the operations accessible to the tellers. Analysts can access statistics for all branches but not perform operations at the account level.

Objects Actors	Corporation	LocalBranch	Account
Teller		lookupLocalAccount ()	postSmallDebit () postSmallCredit () lookupBalance ()
Manager		lookupLocalAccount ()	postSmallDebit () postSmallCredit () postLargeDebit () postLargeCredit () examineBalance () examineHistory ()
Analyst	examineGlobalDebits () examineGlobalCredits ()	examineLocalDebits () examineLocalCredits ()	

- We can represent the access matrix using one of three different approaches: global access table, access control list, and capabilities:
- A **global access table** represents explicitly every cell in the matrix as a (actor, class, operation) tuple. Determining if an actor has access to a specific object requires looking up the corresponding tuple. If no such tuple is found, access is denied.
 - An **access control list** associates a list of (actor, operation) pairs with each classto be accessed. Empty cells are discarded. Every time an object is accessed, its access list is checked for the corresponding actor and operation. An example of an access control list is the guest list for a party. A butler checks the arriving guests by comparing their names against names on the guest list. If there is a match, the guests can enter; otherwise, they are turned away.
 - A **capability** associates a (class, operation) pair with an actor. A capability provides an actor to gain control access to an object of the class described in the capability. Denying a capability is equivalent to denying access. An example of a capability is an invitation card for a party. In this case, the butler checks if the arriving guests hold an invitation for the party. If the invitation is valid, the guests are admitted; otherwise, they are turned away. No other checks are necessary.







The representation of the access matrix is also a performance issue. Global access tables are rarely used as they require a lot of space. Access control lists make it faster to answer the question, “Who has access to this object?” whereas capability lists make it faster to answer the question, “Which objects has this actor access to?”

Each row in the access matrix represents a different access view of the classes listed in the columns. All of these access views should be consistent. Usually, however, access views are implemented by defining a subclass for each different type of (actor, operation) tuple. For example, in our banking system, we would implement an `AccountViewedByTeller` and `AccountViewedByManager` class as subclasses of `Account`. Only the appropriate classes are available to the corresponding actor. For example, the Analyst client software would not include an `Account` class, because the Analyst has no access to any operation in this class. This reduces the risk that an error in the system results in the possibility of unauthorized access.

An access matrix only represents **static access control**. This means that access rights can be modeled as attributes of the objects of the system. In the bank information system example, consider a broker actor who is assigned dynamically a set of portfolios. By policy, a broker cannot access portfolios managed by another broker. In this case, we need to model access rights dynamically in the system, and, hence, this type of access is called **dynamic access control**. For example, Figure 6-38 shows how this access can be implemented with a protection **Proxy pattern** [Gamma et al., 1994]. For each Portfolio, we create a `PortfolioProxy` to protect the Portfolio and check for access. An Access association between a legitimate Broker and a `PortfolioProxy` indicates which Portfolio the Broker has access to. To access a Portfolio, the Broker sends a message to the corresponding `PortfolioProxy`. The `PortfolioProxy` first checks if the invoking Broker has the appropriate association with the `PortfolioProxy`. If access is granted, the `PortfolioProxy` delegates the message to the Portfolio. Otherwise, the operation fails.

In both types of access control, we assume that we know the actor: either the user behind the keyboard or the calling subsystem. This process

of verifying the association between the identity of the user or subsystem and the system is called **authentication**. A widespread authentication mechanism, for example, is for the user to specify a user name, known by everybody, and a corresponding password, only known to the system and stored in an access control list. The system protects its users' passwords by encrypting them before storing or transmitting them. If only a single user knows this user name/password combination, then we can assume that the user behind the keyboard is legitimate. Although password authentication can be made secure with current technology, it suffers from many usability disadvantages: Users choose passwords that are easy to remember and, thus, easy to guess. They also tend to write their password on notes that they keep close to their monitor, and thus, visible to many other, unauthorized, users. Fortunately, other, more secure, authentication mechanisms are available. For example, a smart card can be used in conjunction with a password: An intruder would need both the smart card and the password to gain access to the system. Better, we can use a biometric sensor for analyzing patterns of blood vessels in a person's fingers or eyes. An intruder would then need the physical presence of the legitimate user to gain access to the system, which is much more difficult than just stealing a smart card.










Figure 6-38 Dynamic access implemented with a protection Proxy. The `Access` association class contains a set of operations that `Broker` can use to access a `Portfolio`. Every operation in the `PortfolioProxy` first checks with the `isAccessible()` operation if the invoking `Broker` has legitimate access. Once access has been granted, `PortfolioProxy` delegates the operation to the actual `Portfolio` object. One `Access` association can be used to control access to many `Portfolios`.



In an environment where resources are shared among multiple users, authentication is usually not sufficient. In the case of a network, for example, it is relatively easy for an intruder to find tools to snoop the network traffic, including packets generated by other users (see Figure 6-39). Worse, protocols such as TCP/IP were not designed with security in mind: An intruder can forge packets such that they appear as if they were coming from legitimate users.

Encryption is used to prevent such unauthorized accesses. Using an encryption algorithm, we can translate a message, called **plaintext**, into a encrypted message, called a **ciphertext**, such that even if an intruder intercepts the message, it cannot be understood. Only the receiver has sufficient knowledge to correctly decrypt the message, that is, for reversing the original process. The encryption process is parameterized by a **key**, such that the method of encryption and decryption can be switched quickly in case the intruder manages to obtain sufficient knowledge to decrypt the message.








Figure 6-39 Passive attack. Given current technology, it is relatively easy for a passive intruder to listen to all network traffic. To prevent this kind of attack, encryption makes the information an intruder sees difficult to understand.

Secure authentication and encryption are fundamentally difficult problems. You should always select one or more off-the-shelf algorithms or packages instead of designing your own (unless you are in the business of building such packages). Many such packages are based on public standards that are widely reviewed by the academia and the industry, thus ensuring a relatively high level of reliability and security.

Encapsulating access control

The use of vendor-supplied software introduces a security problem: How can we be sure that the supplied software does not include a trap door? Moreover, once a vulnerability is found in a widely used package, how do we protect the system until a patch is available? We can use redundancy to address both issues. For example, the Java Cryptographic Architecture [JCA, 1998] allows multiple implementations of the same algorithms to coexist in the same system, thus reducing the dependency on a specific vendor. More generally, we can use the **Strategy pattern** [Gamma et al., 1994] to encapsulate multiple implementation of the same algorithm. In this pattern (see Figure 6-40), the *Strategy* abstract class defines the generic interface that all implementations of the encapsulated algorithm should have. Concrete *Strategy* classes provide implementations of the algorithm by subclassing *Strategy*. A *Context* class is responsible for managing the data structure on which Concrete *Strategies* operate. Context and a Concrete *Strategy* class cooperate to provide the needed functionality.

Once authentication and encryption are provided, application specific-access control can be more easily implemented on top of these building blocks. In all cases, addressing security issues is a difficult topic. When addressing these issues, developers should record their assumptions and describe the intruder scenarios they are considering. When several alternatives are explored, developers should state the design problems they are attempting to solve and record the results of the evaluation. We describe in the next chapter how to do this systematically using issue modeling.








Figure 6-40 An example of Strategy pattern encapsulating multiple implementation of the IDEA encryption algorithm (UML class diagram). The *Message* and *IDEA* classes cooperate to realize the encryption of plain text. The selection of an implementation can be done dynamically.







6.4.7 Designing the global control flow

Control flow is the sequencing of actions in a system. In object-oriented systems, sequencing actions includes deciding which operations should be executed and in which order. These decisions are based on external events generated by an actor or on the passage of time.

Control flow is a design problem. During analysis, control flow is not an issue, because we simply assume that all objects are running simultaneously, executing operations any time they need to execute them. During system design, we need to take into account that not every object has the luxury of running on its own processor. There are three possible control flow mechanisms:

- **Procedure-driven control.** Operations wait for input whenever they need data from an actor. This kind of control flow is mostly used in legacy systems and systems written in procedural languages. It introduces difficulties when used with object-oriented languages. As the sequencing of operations is distributed among a large set of objects, it becomes increasingly difficult to determine the order of inputs by looking at the code (Figure 6-41).





```
Stream in, out;  
String userid, passwd;  
/* Initialization omitted */ out.println("Login:"); in.readLine(userid); out.println("Password:"); in.readLine(passwd); if (!security.check(userid, passwd))  
    { out.println("Login failed."); system.exit(-1); }  
/* ... */
```







Figure 6-41 An example of procedure driven control (Java). The code prints out messages and waits for input from the user.

- **Event-driven control.** A main loop waits for an external event. Whenever an event becomes available, it is dispatched to the appropriate object, based on information associated with the event. This kind of control flow has the advantage of leading to a simpler structure and to centralizing all input in the main loop. However, it makes the implementation of multistep sequences more difficult to implement (Figure 6-42).
- **Threads.** Also called light weight threads to distinguish them from processes that require more computing overhead, are the concurrent variation of procedure-driven control: The system can create an arbitrary number of threads, each responding to a different event. If a thread needs additional data, it waits for input from a specific actor. This kind of control flow is probably the most intuitive of the three mechanisms. However, debugging threaded software requires good debugging tools: Preemptive thread schedulers introduce nondeterminism and, thus, make it harder to find repeatable test cases (Figure 6-43).



```
Enumeration subscribers, eventStream; Subscriber subscriber; Event event; EventStream eventStream;  
/* ... */  
while (eventStream.hasMoreElements()) { event = eventStream.nextElement(); subscribers = dispatchInfo.getSubscribers(event); while (subscribers.hasMoreElements()) {  
    subscriber = subscribers.nextElement(); subscriber.process(event); } }  
/* ... */
```

Figure 6-42 An example of main loop for event-driven control (Java). An event is taken from an event queue and sent to objects interested in it.





```
Thread thread; Event event; EventHandler eventHandler; boolean done;
/* ... */
while (!done) { event = eventStream.getNextEvent(); eventHandler = new EventHandler(event) thread = new Thread(eventHandler); thread.start();
}
/* ... */
```




Figure 6-43 An example of event processing with threads (Java). eventHandler is an object dedicated to handling event. It implements the run() operation, which is invoked when thread is started.




Procedure-driven control is useful for testing subsystems. A driver makes specific calls to methods offered by the subsystem. For the control flow of the final system, though, procedure-driven control should be avoided.

The trade-off between event-driven control and threads is more complicated. Event-driven control is more mature than threads. Modern languages have only recently started to provide support for thread programming. As more debugging tools are becoming available and experience is accumulated, developing thread-based systems will become easier. Also, many user interface packages supply the infrastructure for dispatching events and impose this kind of control flow on the design. Although threads are more intuitive, they currently introduce many problems during debugging and testing. Until more mature tools and infrastructures are available for developing with threads, event-driven control flow is preferred.

Once a control flow mechanism is selected, we can realize it with a set of one or more control objects. The role of control objects is to record external events, store temporary state about them, and issue the right sequence of operation calls on the boundary and entity objects associated with the external event. On the one hand, localizing control flow decisions for a use case into a single object results into more understandable code; on the other hand, it makes the system more resilient to changes in control flow implementation.

Encapsulating control flow

An example of encapsulation of control is the **Command pattern** [Gamma et al., 1994] (see Figure 6-44). In interactive systems, it is often desirable to execute, undo, or store user requests without knowing the content of the request. The key to decoupling requests from their handling is to turn requests into command objects, which inherit from an abstract *Commandclass*. The *Commandclass* defines how the command is executed, undone, or stored, whereas the concrete class implements specific requests.






Figure 6-44 Command pattern (UML class diagram). This pattern enables the encapsulation of control such that user requests can be treated uniformly, independent of the specific request.



We can use the Command pattern to decouple menu items from actions (see Figure 6-45). Decoupling menu items from actions has the advantage of centralizing control flow (e.g., dialog sequencing) into control objects instead of spreading it between boundary and entity objects. A Menu, composed of MenuItem s, creates a Command object of the appropriate class whenever the corresponding MenuItem is selected by the user. The Application invokes the execute() operation of the newly created Command object. If the user wishes to undo the last request, the undo() operation of the last Command object is executed. Different Command objects implement different requests (e.g., CopyCommand and PasteCommand).






Figure 6-45 An example of a Command pattern (UML class diagram). In this example, menu items and operations on documents are decoupled. This enables us to centralize control flow in the command objects (CopyCommand and PasteCommand) instead of spreading it across boundary objects (MenuItem) and entity objects (Document).





6.4.8 Identifying boundary conditions

In previous sections, we dealt with designing and refining the system decomposition. We now have a better idea of how to decompose the system, how to distribute use cases among subsystems, where to store data, and how to achieve access control and ensure security. We still need to examine the boundary conditions of the system, that is, to decide how the system is started, initialized, and shut down, and we need to define how we deal with major failures, such as data corruption, whether they are caused by a software error or a power outage.






Figure 6-46 Administration use cases for MyTrip (UML use case diagram). *ManageDrivers* is invoked to add, remove, modify, or read data about drivers (e.g., user name and password, usage log, encryption key generation). *ManageMaps* is invoked to add, remove, or update maps that are used to generate trips. *ManageServer* includes all the functions necessary to start up and shutdown the server.



For example, we now have a good idea of how MyTrip should work in steady state. We have, however, not yet addressed how MyTrip is initialized. For example, how are maps loaded into the PlanningService? How is MyTrip installed in the car? How does MyTrip know which PlanningService to connect to? How are drivers added to the PlanningService? We quickly discover a set of use cases that has not been specified. We call these the system administration use cases. *System administration use cases* specify the behavior of a system during the start-up and shutdown phase.

It is common that system administration use cases are not specified during analysis or that they are treated separately. On the one hand, many system administration functions can be inferred from the everyday user requirements (e.g., registering and deleting users, managing access control); on the other hand, many functions are consequences of design decisions (e.g., cache sizes, location of database server, location of backup server) and not of requirement decisions.





We now modify the analysis model for MyTrip to include the administration use cases. In particular, we add three use cases: ManageDrivers, to add, remove, and edit drivers; ManageMaps, to add, remove, and update maps used to generate trips; and ManageServer, to perform routine configuration, start-up, and shutdown of the server (see Figure 6-46). StartServer, part of ManageServer, is provided as an example in Figure 6-47.

case	
name	
Entry	1. The PlanningServiceAdministrator logs into the server machine.
conditi	
on	
Flow	2. Upon successful login, the PlanningServiceAdministrator executes the startPlanningService command.
of	
events	3. If the PlanningService was previously shutdown normally, the server reads the list of legitimate Drivers and the index of active Trips and Maps. If the PlanningService had


In this case, adding three use cases, that is, revising the use case model, does not impact the subsystem decomposition. We added, however, new use cases to existing subsystems: The MapDBStoreSubsystem needs to be able to detect whether it was properly shut down or not and needs to be able to perform consistency checks and repair corrupted data, if necessary. We revise the description of MapDBStoreSubsystem (Figure 6-48).



When examining boundary conditions, we also need to investigate exceptional cases. For example, the nonfunctional requirements of MyTrip specify that the system needs to tolerate connection failures. For this reason, the RouteAssistant downloads the planned Trip at the initial Destination. We also decide to download segments that are close to the Trip, to enable minimum replanning even though a connection might not be available.



MapDBStoreSubsystem The MapDBStoreSubsystem is responsible for storing maps and trips in a database for the PlanningSubsystem. This subsystem supports multiple concurrent drivers and planning agents. *When starting up, the MapDBStoreSubsystem detects if it was properly shutdown. If not, it performs a consistent check on the Maps and Trips and repairs corrupted data if necessary.*

Figure 6-48 Revised description for MapDBStoreSubsystem based on the additional StartServer use case of Figure 6-47. (Changes indicated in *italics*.)





In general, an **exception** is an unexpected event or error that occurs during the execution of the system. Exceptions are caused by one of three different sources:

- *A user error.* The user mistakenly or deliberately inputs data that are out of bounds. For example, a negative amount in a banking transaction could lead to transferring money in the wrong direction if the system does not protect against such errors.
- *A hardware failure.* Hardware ages and fails. The failure of a network link, for example, can momentarily disconnect two nodes of the system. A hard disk crash can lead to the permanent loss of data.
- *A software bug.* An error can occur either because the system or one of its components contains a design error. Although writing bug-free software is difficult, individual subsystems can anticipate errors from other subsystems and protect against them.




Exception handling is the mechanism by which a system treats an exception. In the case of a user error, the system should display a meaningful error message to the user, such that she can correct her input. In the case of a network link failure, the system needs to save its temporary state in order to recover once the network comes back on line.


Developing reliable systems is a difficult topic. Often, trading off some functionality can make it easier on the design of the system. In MyTrip, we assumed that the connection is always possible at the source destination and that replanning could be impacted by communication problems along the trip.

6.4.9 Anticipating change

System design introduces a strange paradox in the development process. On the one hand, we want to construct solid walls between subsystems, to manage complexity by breaking the system into smaller pieces, and to prevent changes from one subsystem to impact others. On the other hand, we want the software architecture to be modifiable to minimize the cost of later change. These are conflicting goals that cannot be reconciled: We have to define an architecture early to deal with complexity, and we have to pay the price of change later in the development process. We can, however, anticipate change and design for it, as sources of later changes tend to be the same for most systems:

- *New vendor or new technology.* When components are used to build the system, anticipate that the component will be replaced by an equivalent one from a different vendor. This change is frequent and generally difficult to cope with. The software marketplace is dynamic, and many vendors will start up and go out of business before your project is completed.
- *New implementation.* When subsystems are integrated and tested together, the overall system response time is, more often than not, above stated or implicit performance requirements: Posting a debit on a bank information system may take 2 minutes, a flight reservation system takes 5 minutes to book a flight. System-wide performance is difficult to predict and is usually not optimized before integration: Developers focus on their subsystem first. This triggers the need for more-efficient data structures and algorithms and better interfaces—often under time constraints.
- *New views.* Testing the software with real users uncovers many usability problems. These often translate into the creation of additional views on the same data.
- *New complexity of the application domain.* The deployment of a system triggers ideas of new generalizations: A bank information system for one branch may lead to the idea of a multibranch information system. Other times, the domain itself increases in complexity: Previously, flight numbers were associated with one plane and one plane only. With the advent of carrier alliances, one plane can now have multiple flight numbers from different companies.
- *Errors.* Unfortunately, many requirements errors are discovered only when real users start using the system.





Adapter
(see
example in
Figure 6-
34)

New vendor, new technology, new implementation. This pattern encapsulates a piece of legacy code that was not designed to work with the system. It also limits the impact of substituting the piece of legacy code for a different component.

Bridge
(see
example in
Figure 6-
37)






New vendor, new technology, new implementation. This pattern decouples the interface of a class from its implementation. It serves the same purpose as the Adapter pattern except that the developer is not constrained by an existing piece of code.

Command
(see
example in
Figure 6-
45)


New functionality. This pattern decouples the objects responsible for command processing from the commands themselves. This pattern protects these objects from changes due to new functionality.

Observer
(see
example in
Section
6.3.5)

New views. This pattern decouples entity objects from their views. Additional views can be added with entity objects being modified.



Modern object-oriented languages provide mechanisms that can minimize the impact of change when anticipated. The use of delegation and inheritance in conjunction with abstract classes decouples the interface of a subsystem from its actual implementation. In this chapter, we have provided you with selected examples of design patterns [Gamma et al., 1994] that deal with the above changes. Figure 6-49 summarizes the patterns and the type of change they protect against.



A reason for the high cost of change late in the process is the loss of design context. Developers forget very quickly the reasons that pushed them to design complicated workarounds or complex data structures during early phases of the process. When changing code late in the process, the likelihood of introducing errors into the system is high. To protect against such situations, assumptions should be recorded. For example, when using a design pattern to anticipate a certain change (from Figure 6-49), you should record which change you are anticipating. In Chapter 8, *Rationale Management*, we describe several techniques for recording the design alternatives and decisions.

6.4.10 Reviewing system design

Like analysis, system design is an evolutionary and iterative activity. Unlike analysis, there is no external agent, such as the client, to review the successive iterations and ensure better quality. This quality improvement activity, however, is still necessary, and project managers and developers need to organize a review process to substitute for it. Several alternatives exist, such as using the developers who were not involved in system design to act as independent reviewers, or to use developers from another project to act as a peer review. These review processes work only if the reviewers have an incentive in discovering and reporting problems.




In addition to meeting the design goals that were identified during system design, we need to ensure that the system design model is correct, complete, consistent, realistic, and readable. The system design model is **correct** if the analysis model can be mapped to the system design model. You should ask the following questions to determine if the system design is correct:


- Can every subsystem be traced back to a use case or a nonfunctional requirement?

- Can every use case be mapped to a set of subsystems?
- Can every design goal be traced back to a nonfunctional requirement?
- Is every nonfunctional requirement addressed in the system design model?
- Does each actor have an access policy?
- Is it consistent with the nonfunctional security requirement?

The model is **complete** if every requirement and every system design issue has been addressed. You should ask the following questions to determine if the system design is complete:

- Have the boundary conditions been handled?
- Was there a walkthrough of the use cases to identify missing functionality in the system design?
- Have all use cases been examined and assigned a control object?
- Have all aspects of system design (i.e., hardware allocation, persistent storage, access control, legacy code, boundary conditions) been addressed?
- Do all subsystems have definitions?





The model is **consistent** if it does not contain any contradictions. You should ask the following questions to determine if a system design is consistent:

- Are conflicting design goals prioritized?
- Are there design goals that violate a nonfunctional requirement?
- Are there multiple subsystems or classes with the same name?
- Are collections of objects exchanged among subsystems in a consistent manner?

The model is **realistic** if the corresponding system can be implemented. You ask the following questions to determine if a system design is realistic:

- Are there any new technologies or components in the system? Were there any studies to evaluate the appropriateness or robustness of these technologies or components?
- Have performance and reliability requirements been reviewed in the context of the subsystem decomposition? For example, is there a network connection on the critical path of the system?
- Have concurrency issues (e.g., contention, deadlocks, mutual exclusion) been addressed?

The model is **readable** if developers not involved in the system design can understand the model. You should ask the following questions to ensure that the system design is readable:




- Are subsystem names understandable?
- Do entities (e.g., subsystems, classes, operations) with similar names denote similar phenomena?
- Are all entities described at the same level of detail?

In many projects, you will find that system design and implementation overlap quite a bit. For example, you may build prototypes of selected subsystems before the architecture is stable in order to evaluate new technologies. This leads to many partial reviews instead of an encompassing review followed by a client sign-off, as for analysis. Although this process yields greater flexibility, it also requires developers to track open issues more carefully. Many difficult issues tend to be resolved late, not because they are difficult, but because they fell through the cracks of the organization.

6.5 Managing system design

In this section, we discuss issues related to managing the system design activities. As in analysis, the primary challenge in managing the system design is to maintain consistency, while using as many resources as possible. In the end, the software architecture and the system interfaces should describe a single coherent system understandable by a single person.


We first describe a document template that can be used to document the results of system (Section 6.5.1). Next, we describe the role assignment during system design (Section 6.5.2) and address communication issues during system design (Section 6.5.3). Next, we address management issues related to the iterative nature of system design (Section 6.5.4).





6.5.1 Documenting system design


System design is documented in the System Design Document (SDD). It describes design goals set by the project, the subsystem decomposition (with UML class diagrams), the hardware/software mapping (with UML deployment diagrams), the data management, the access control, control flow mechanisms, and boundary conditions. The SDD is used to define interfaces between teams of developers and as reference when architecture-level decisions need to be revisited. The audience for the SDD includes the project management, the system architects (i.e., the developers who participate in the system design), and the developers who design and implement each subsystem. The following is an example template for a SDD:



System Design Document


- 1. Introduction
 - 1.1 Purpose of the system
 - 1.2 Design goals
 - 1.3 Definitions, acronyms, and abbreviations
 - 1.4 References
 - 1.5 Overview
- 2. Current software architecture
- 3. Proposed software architecture



Overview
Subsystem decomposition
Hardware/software mapping
Persistent data management
Access control and security
Global software control
Boundary conditions
Subsystem services Glossary



The first section of the SDD is an *Introduction*. Its purpose is to provide a brief overview of the software architecture and the design goals. It also provides references to other documents and traceability information (e.g., related requirements analysis document, references to existing systems, constraints impacting the software architecture).

The second section, *Current software architecture* describes the architecture of the system being replaced. If there is no previous system, this section can be replaced by a survey of current architectures for similar systems. The purpose of this section is to make explicit the background








information that system architects used, their assumptions, and common issues the new system will address.



The third section, *Proposed system architecture* documents the system design model of the new system. It is divided into seven subsections:

- *Overview* presents a bird's eye view of the software architecture and briefly describes the assignment of functionality to each subsystem.
- *Subsystem decomposition* describes the decomposition into subsystems and the responsibilities of each. This is the main product of system design.
- *Hardware/software mapping* describes how subsystems are assigned to hardware and off-the-shelf components. It also lists the issues introduced by multiple nodes and software reuse.
- *Persistent data management* describes the persistent data stored by the system and the data management infrastructure required for it. This section typically includes the description of data schemes, the selection of a database, and the description of the encapsulation of the database.
- *Access control and security* describes the user model of the system in terms of an access matrix. This section also describes security issues, such as the selection of an authentication mechanism, the use of encryption, and the management of keys.
- *Global software control* describes how the global software control is implemented. In particular, this section should describe how requests are initiated and how subsystems synchronize. This section should list and address synchronization and concurrency issues.
- *Boundary conditions* describes the start-up, shutdown, and error behavior of the system. If new use cases are discovered for system administration, these should be included in the requirements analysis document, not in this section.

The fourth section, *Subsystem services*, describes the services provided by each subsystem in terms of operations. Although this section is usually empty or incomplete in the first versions of the SDD, this section serves as a reference for teams for the boundaries between their subsystems. The interface of each subsystem is derived from this section and detailed in the Object Design Document.

The SDD is written after the initial system decomposition is done; that is, system architects should not wait until all system design decisions are made before publishing the document. The SDD, moreover, is updated throughout the process when design decisions are made or problems are discovered. The SDD, once published, is baselined and put under configuration management. The revision history section of the SDD provides a history of changes as a list of changes, including author responsible for the change, date of change, and brief description of the change.





6.5.2 Assigning responsibilities

Unlike analysis, system design is the realm of developers. The client and the end user fade into the background. Note, however, that many activities in system design trigger revisions to the analysis model. The client and the user are brought back into the process for such revisions. System design in complex systems is centered around the architecture team. This is a cross-functional team made up of architects (who define the subsystem decomposition) and selected developers (who will take part in the implementation of the subsystem). It is critical that system design include persons that are exposed to the consequences of system design decisions. The architecture team starts work right after the analysis model is stable and continue to function until the end of the integration phase. This creates an incentive for the architecture team to anticipate problems encountered during integration. Below are the main roles of system design:




- The **architect** is the main role of system design. The architect ensures consistency in design decisions and interface styles. The architect ensures the consistency of the design in the configuration management and testing teams, in particular in the formulation of the configuration management policy as well as the system integration strategy. This is mainly an integration role consuming information from each subsystem team. The architect is the leader of the cross-functional architecture team.
- **Architecture liaisons** are the members of the architecture team. They are representatives from the subsystem teams. They convey information from and to their teams and negotiate interface changes. During system design, they focus on the subsystem services; during the implementation phase, they focus on the consistency of the APIs.
- The **document editor**, **configuration manager**, and **reviewer** roles are the same as for analysis.



The number of subsystems determines the size of the architecture team. For complex systems, an architecture team is introduced for each level of abstraction. In all cases, there should be one integrating role on the team to ensure consistency and the understandability of the architecture by a single individual.

6.5.3 Communicating about system design

Communication during system design should be less challenging than during analysis: The functionality of the system has been defined, project participants have similar backgrounds and by now should know each other better. Communication is still difficult, due to new sources of complexity:




- *Size.* The number of issues to be dealt with increases as developers start designing. The number of items that developers manipulate increases: Each piece of functionality requires many operations on many objects. Moreover, developers investigate, often concurrently, multiple designs and multiple implementation technologies.
- *Change.* The subsystem decomposition and the interfaces of the subsystems are in constant flux. Terms used by developers to name different parts of the system evolve constantly. If the change is rapid, developers may not be discussing the same version of the subsystem, which can lead to much confusion.
- *Level of abstraction.* Discussions about requirements can be made concrete by using interface mock-ups and analogies with existing systems. Discussions about implementation become concrete when integration and test results are available. System design discussions are seldom concrete, as consequences of design decisions are felt only later, during implementation and testing.
- *Reluctance to confront problems.* The level of abstraction of most discussions can also make it easy to delay the resolution of difficult issues. A typical resolution of control issues is often, “Let us revisit this issue during implementation.” Whereas it is usually desirable to delay certain design decisions, such as the internal data structures and algorithms used by each subsystem, for example, any decision that has an impact on the system decomposition and the subsystem interfaces should not be delayed.
- *Conflicting goals and criteria.* Individual developers often optimize different criteria. A developer experienced in user interface design will be biased toward optimizing response time. A developer experienced in databases might optimize throughput. These conflicting goals, especially when implicit, result in developers pulling the system decomposition in different directions and lead to inconsistencies.







The same techniques we discussed in analysis (see Section 5.5.3) can be applied during system design:

- *Identify and prioritize the design goals for the system and make them explicit* (see Section 6.4.2). If the developers concerned with system design have an input in this process, they will have an easier time committing to these design goals. Design goals also provide an objective framework against which decisions can be evaluated.
- *Make the current version of the system decomposition available to all concerned.* A live document distributed via the Internet is one way to achieve rapid distribution. Using a configuration management tool to maintain the system design documents helps developers in identifying recent changes.
- *Maintain an up-to-date glossary.* As in analysis, defining terms explicitly reduces misunderstandings. When identifying and modeling subsystems, provide definitions in addition to names. A UML diagram with only subsystem names is not sufficient for supporting effective communication. A brief and substantial definition should accompany every subsystem and class name.
- *Confront design problems.* Delaying design decisions can be beneficial when more information is needed before committing to the design decision. This approach, however, can prevent the confrontation of difficult design problems. Before tabling an issue, several possible alternatives should be explored and described and the delay justified. This ensures that issues that are delayed can be delayed without serious impact on the system decomposition.
- *Iterate.* Selected excursions into the implementation phase can improve the system design. For example, new features in a vendor-supplied component can be evaluated by implementing a vertical prototype (see Section 6.5.4) for the functionality most likely to benefit from the feature.





Finally, no matter how much effort is expended on system design, the system decomposition and the subsystem interfaces will almost certainly change during implementation. As new information about implementation technologies becomes available, developers have a clearer understanding of the system, and design alternatives are discovered. Developers should anticipate change and reserve some time to update the SDD before system integration.




6.5.4 Iterating over the system design



As in the case of requirements, system design occurs through successive iteration and change. Change, however, should be controlled to prevent chaos, especially in complex projects including many participants. We distinguish three types of iterations during system design. First, major decisions early in the system design impact subsystem decomposition as each of the different activities of system design are initiated. Second, revisions to the interfaces of the subsystems occur when evaluation prototypes are done to evaluate specific issues. Third, errors and oversights that are discovered late trigger changes to the subsystem interfaces and sometimes to the system decomposition itself.

The first set of iterations is best handled through face-to-face and electronic brainstorming. Definitions are still in flux, developers do not have yet a grasp of the whole system, and communication should be maximized at the expense of formality or procedure. Often in team-based projects, the initial system decomposition is designed before the analysis is complete. Decomposing the system early enables the responsibility of different subsystems to be assigned to different teams. Change and exploration should be encouraged, if only to broaden the developers shared understanding or to generate supporting evidence for the current design. For this reason, there should not be a bureaucratic formal change process during this

phase.

The second set of iterations aims at solving difficult and focused issues, such as the choice of a specific vendor or technology. The subsystem decomposition is stable (it should be independent of vendors and technology; see Section 6.4.9), and most of these explorations aim at identifying whether a specific package is appropriate or not for the system. During this period, developers can also realize a vertical prototype⁴ for a critical use case to test the appropriateness of the decomposition. This enables control flow issues to be discovered early and addressed. Again, a formal change process is not necessary. A list of pending issues and their status can help developers quickly propagate the result of a technology investigation.





The third set of iterations remedies design problems discovered late in the process. Although developers would much rather avoid any such iterations, as they tend to incur a high cost and introduce many new bugs in the system, they should anticipate changes late in the development. Anticipating late iterations includes documenting dependencies among subsystems, the design rationale for subsystem interfaces, and any workaround that is likely to break in case of change. Change should be carefully managed, and a change process, similar to the one tracking requirements changes, should be put in place.

We can achieve the progressive stabilization of the subsystem decomposition by using the concept of design window. In order to encourage change while controlling it, critical issues are left open only during a specified time. For example, the hardware/software platform on which the system is targeted should be resolved early in the project, such that purchasing decisions for the hardware can be done in time for the developers. Internal data structures and algorithms, however, can be left open until after integration, allowing developers to revise them based on performance testing. Once the design window is past, the issue must be resolved and can only be reopened in a subsequent iteration.

With the pace of technology innovation quickening, many changes can be anticipated when a dedicated part of the organization is responsible for technology management. Technology managers scan new technologies, evaluate them, and accumulate knowledge that is used during the selection of components. Often, change happens so fast that companies are not aware of which technologies they themselves provide.

6.6 Exercises





1. Decomposing a system into subsystems reduces the complexity developers have to deal with by simplifying the parts and increasing their


coherence. Decomposing a system into simpler parts usually results into increasing a different kind of complexity: Simpler parts also means a larger number of parts and interfaces. If coherence is the guiding principle driving developers to decompose a system into small parts, which competing principle drives them to keep the total number of parts small?

2. In Section 6.4.2, we classified design goals into five categories: performance, dependability, cost, maintenance, and end user. Assign one or more categories to each of the following goals:

- Users must be given a feedback within 1 second after they issue any command.
- The TicketDistributormust be able to issue train tickets, even in the event of a network failure.
- The housing of the TicketDistributormust allow for new buttons to be installed

4. A vertical prototype completely implements a restricted functionality (e.g., interface, control, and entity objects for one use case), whereas a horizontal prototype implements partially a broad range of functionality (e.g., interface objects for a number of use cases).







in the event the number of different fares increases.



- The AutomatedTellerMachinemust withstand dictionary attacks (i.e., users attempting to discover a identification number by systematic trial).
 - The user interface of the system should prevent users from issuing commands in the wrong order.
1. Consider a system that includes a Web server and two database servers. Both database servers are identical: The first acts as a main server, while the second acts as a redundant backup in case the first one fails. Users use Web browsers to access data through the Web server. They also have the option of using a proprietary client that accesses the databases directly. Draw a UML deployment diagram representing the hardware/software mapping of this system.
 2. Consider a legacy, fax-based, problem-reporting system for an aircraft manufacturer. You are part of a reengineering project replacing the core of the system by a computer-based system, which includes a database and a notification system. The client requires the fax to remain an entry point for problem reports. You propose an E-mail entry point. Describe a subsystem decomposition, and possibly a design pattern, which would allow both interfaces.
 3. You are designing the access control policies for a Web-based retail store. Customers access the store via the Web, browse product information, input their address and payment information, and purchase products. Suppliers can add new products, update product information, and receive orders.The store owner sets the retail prices, makes tailored offers to customers based on their purchasing profiles, and provides marketing services. You have to deal with three actors: StoreAdministrator. Design

, Supplier, and Customer access control policy for all three actors. Customers can be created via the Web\, whereas Suppliers are created by the StoreAdministrator

6. Select a control flow mechanism you find most appropriate for each of the following systems. Because multiple choices are possible in most cases, justify your choices.

- a Web server designed to sustain high loads
- a graphical user interface for a word processor
- a real-time embedded system (e.g., a guidance system on a satellite launcher)
- Why are use cases that describe boundary conditions described during system design (as opposed to during requirements elicitation or analysis)?
- You are developing a system that stores its data on a Unix file system. You anticipate that you will port future versions of the system to other operating systems that provide different file systems. Which design pattern do you use to guard against this change?





References

[Bass et al., 1999] [Booch, 1994] [Day &
Zimmermann, 1983] [Erman et al., 1980]

[Fowler, 1997]

[Gamma et al., 1994]

[JCA, 1998] [JDBC, 1998]

[Microsoft, 1995]

[Mowbray & Malveau, 1997] [Nye & O'Reilly, 1992]

[OMG, 1995]

[RMI, 1998] [Rumbaugh et al., 1991] [Shaw & Garlan, 1996] [Siewiorek & Swarz, 1992] [Silberschatz et al, 1991] [Tanenbaum,
1996]

L. Bass, P. Clements, & R. Kazman, *Software Architecture in Practice*. Addison-Wesley, Reading, MA, 1999.

G. Booch, *Object-Oriented Analysis and Design with Applications*, 2nd ed. Benjamin/Cummings, Redwood City, CA, 1994.

- Day & H. Zimmermann, “The OSI Reference Model,” *Proceedings of the IEEE*, vol. 71, pp. 1334–1340, Dec. 1983.
- D. Erman, F. Hayes-Roth, et al., “The Hearsay-II Speech-Understanding System: Integrating knowledge to resolve uncertainty.” *ACM Computing Surveys*, vol. 12, no. 2, pp. 213–253, 1980.
- Fowler, *Analysis Patterns: Reusable Object Models*. Addison-Wesley, Reading, MA, 1997.

• Gamma, R. Helm, R. Johnson, & J. Vlissides, *Design Patterns: Elements of Reusable Object-Oriented Software*. Addison-Wesley, Reading, MA, 1994.

Java Cryptography Architecture, JDK Documentation. Javasoft, 1998.

JDBC™—Connecting Java and Databases, JDK Documentation. Javasoft, 1998.

Microsoft, “Chapter 9: Open Database Connectivity (ODBC) 2.0 fundamentals,” *Microsoft Windows Operating Systems and Services Architecture*, Microsoft Corp., 1995.

T. J. Mowbray & R. C. Malveau, *CORBA Design Patterns*. Wiley, New York, 1997.

A. Nye & T. O’Reilly. *X Toolkit Intrinsic Programming Manual: OSF/Motif 1.1 Edition for X11 Release 5 The Definitive Guides to the X Windows Systems*, vol. 4, O’Reilly & Associates, Sebastopol, CA, 1992.

Object Management Group, *The Common Object Request Broker: Architecture and Specification*. Wiley, New York, 1995.

Java Remote Method Invocation, JDK Documentation. Javasoft, 1998.




J. Rumbaugh, M. Blaha, W. Premerlani, F. Eddy, and W. Lorensen. *Object-Oriented Modeling and Design*. Prentice Hall, Englewood Cliffs, NJ, 1991.


M. Shaw & D. Garlan, *Software Architecture: Perspectives on an Emerging Discipline*. Prentice Hall, Upper Saddle River, NJ, 1996.

D. P. Siewiorek & R. S. Swarz, *Reliable Computer Systems: Design and Evaluation*. 2nd ed. Digital, Burlington, MA, 1992.

A. Silberschatz, J. Peterson, & P. Galvin, *Operating System Concepts*. 3rd ed. Addison-Wesley, Reading, MA, 1991.

A. S. Tanenbaum, *Computer Networks*, 3rd ed. Prentice Hall, Upper Saddle River, NJ, 1996.





7.1 Introduction: Movie examples 232

7.2 An overview of object design 233

7.3 Object design concepts 235

7


- 7.3.1 Application objects versus solution objects revisited 237
- 7.3.2 Types, signatures, and visibility revisited 237
- 7.3.3 Contracts: Invariants, preconditions, and postconditions 238
- 7.3.4 UML's Object Constraint Language 240


7.4 Object design activities 241

- 7.4.1 Identifying missing attributes and operations 245
- 7.4.2 Specifying type signatures and visibility 248
- 7.4.3 Specifying constraints 250
- 7.4.4 Specifying exceptions 252
- 7.4.5 Identifying and adjusting class libraries 253
- 7.4.6 Identifying and adjusting application frameworks 255
- 7.4.7 A framework example: WebObjects 257
- 7.4.8 Realizing associations 259
- 7.4.9 Increasing reuse 265
- 7.4.10 Removing implementation dependencies 267
- 7.4.11 Revisiting access paths 270
- 7.4.12 Collapsing objects: Turning objects into attributes 271
- 7.4.13 Caching the result of expensive computations 272
- 7.4.14 Delaying expensive computations 272

7.5 Managing object design 273

- 7.5.1 Documenting object design 274
- 7.5.2 Assigning responsibilities 278





蠱梔最甄熾攀皀 娒袷濟攀矚愀洵漫氈攀漿昀鸛唳糾愀熾振梔梔玲攀振玲甄熾攀「吁梔攀 ㄣ 洵漿揆攀氈 v 梔玳玲梔攀昀梔氈攀濟愀洵攀忌膂攀玳梔最濟
 佹袷玲攀熾濟玳()「漫漫玲「佻濟攀 ㄣ 瘀梔攀明 v 梔玳愀明梔濟揆漿明玲梔玲氈攀揆網漿洵漫 ㄣ 箭愀玳攀揆勻漿昀玲明愀熾攀菱濟最梔濟攀熾梔梔最 卄 明梔梔
 振梔揆梔玳漫氈愀礪玳梔攀振漿濟玲攀濟玲玳漿昀愀昀漿氈揆攀熾振漿濟玲愀梔濟梔濟最玲梔攀昀梔氈攀忌膂攀玳梔最濟佹愀玲攀熾濟玳()「漫漫玲「吁
 梔攀漿玲梔攀熾 ㄣ 瘀梔攀明 v 梔玳明梔濟揆漿明振愀氈氈攀揆忌膂攀玳梔最濟佹愀玲攀熾濟玳()「漫漫玲醋濟昀漿 卄 明梔梔振梔揆梔玳漫氈愀礪玳梔濟昀
 漿熾洵愀玲梔漿濟熾攀氈愀玲攀揆玲漿玲梔攀昀梔氈攀「酷昀玲梔攀昀梔氈攀濟愀洵攀梔玳振梔愀濟最攀揆 卄 戀漿玲梔瘀梔攀明玳愀熾攀甄漫揆愀玲攀
 揆戀礪玲梔攀 ㄣ 振漿濟玲熾漿氈氈攀熾 v「